

ARBITER · SECURITY COMPLIANCE

Security Compliance Coverage Report

Adversarial Test Harness — Verification Matrix

Generated 2026-06-10

71%


FULLY COVERED

56 / 79 items fully covered · **59 / 79** at least partial · **17 / 79** covered by other controls (out of scope for this runtime harness).

This document maps every item on the ARBITER security compliance checklist (OWASP Top 10, API Top 10, LLM Top 10, and related CWE weaknesses) to a concrete test in the adversarial harness at `tests-adversarial/`. For each item it states whether a test exists today, where to find it, and — for partial and missing rows — the gap and the estimated effort to close it. The 17 items marked out of scope are routed to their proper control (SCA, IAM static analysis, incident response runbook, etc.) rather than ignored, so an auditor sees the complete picture.

Source: `docs/security_compliance_coverage.md` · Generated by `scripts/compliance_report.py`

Executive summary



Coverage by category

CATEGORY	COVERED	PARTIAL	MISSING	OOS	TOTAL
1. Injection	10	0	0	0	10
2. Authentication & Access Control	10	1	0	1	12
3. Crypto & Data Handling	3	0	2	1	6
4. Configuration & Infrastructure	5	1	1	3	10
5. Build / Supply Chain	2	0	0	2	4
6. Errors & State	5	0	0	0	5
7. API Top-10 Specific	6	1	0	1	8
8. Client-Side	5	0	0	0	5
9. Logic & Workflow	1	0	0	2	3
10. Denial of Service	2	0	0	1	3
11. Logging & Monitoring	3	0	0	3	6
12. LLM Top-10	4	0	0	2	6
Totals	56	3	3	17	79

What this harness tested

The harness runs nine layers, each a focused pytest or Playwright suite: `e2e`, `fuzz`, `auth`, `headers`, `llm`, `dos`, `logic`, `logging_audit`, `fault`. Each layer's results land in a shared coverage matrix aggregator (`src/coverage/builder.py`) that emits the daily PDF report and the per-run findings JSON. This document is the compliance-side view of that same data: which categories of risk the nine layers exercise, and which they deliberately do not.

Detailed coverage matrix

1. Injection 10 COVERED · 0 PARTIAL · 0 MISSING · 0 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
1	SQL injection	COVERED	`fuzz/corpus/sqli.json` (9 curated payloads) parameterised across all 25 API routes. ARBITER uses DynamoDB so SQLi isn't the literal threat, but the harness catches any string-template logic that mishandles quoted input.
2	NoSQL injection	COVERED	Block A: `fuzz/corpus/nosql_operators.json` (7 entries: `\$ne`, `\$gt`, `\$where`, `\$regex`, `\$or`, etc.) parameterised across all 25 routes.
3	OS command injection	COVERED	`fuzz/corpus/command_injection.json` (6 entries: backticks, `\$( )`, `;`, `
4	LDAP injection	COVERED	Block A: `fuzz/corpus/ldap.json` (6 entries: `*)(uid=*`, `*)(\
5	XPath injection	COVERED	Block A: `fuzz/corpus/xpath.json` (6 entries: `'  or '1'='1`, `or count(/*)>0`, union wildcard).
6	XML injection (incl. XXE)	COVERED	Block A: `fuzz/corpus/xml_xxe.json` (6 entries: file://etc/passwd, billion-laughs, parameter-entity OOB, IMDS-via-XXE).
7	Cross-site scripting (XSS)	COVERED	`fuzz/corpus/xss.json` (14 entries incl. OWASP polyglot, ` <svg +="" <b="" every="" javascript:="" on="" onload>`,="" reflection-detection="" response.="" urls)="">Note: classifier has a 6-char floor to avoid false positives — flagged in earlier fuzz review.</svg>
8	Server-Side Template Injection (SSTI)	COVERED	Block A: `fuzz/corpus/ssti.json` (7 entries: `\${7*7}`, `\${{7*7}}`, `<%= 7*7 %>`, SpEL RCE, Twig dump). Evaluated-form detection (literal `49` in response) is documented as future enhancement.
9	Header / CRLF injection	COVERED	`fuzz/corpus/header_injection.json` (5 entries including CRLF + Set-Cookie smuggling). Probes every route.

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
10	Log injection	COVERED	Block A: `fuzz/corpus/log_injection.json` (5 entries: CR/LF, ANSI ESC[2K erase-line, NUL truncation, ANSI color codes) — API-layer reflection check. Block G: `logging_audit/test_log_injection_downstream.py` adds the CloudWatch verification — for each corpus payload, asserts the payload lands in a single CloudWatch event (no CR/LF split → no log forgery) with no ANSI residue.

2. Authentication & Access Control 10 COVERED · 1 PARTIAL · 0 MISSING · 1 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
11	Insecure Direct Object Reference (IDOR)	COVERED	Block C: `auth/test_idor.py` — CISO creates a session via `POST /chat`; SOC/GRC/Employee each attempt `GET /conversations/{id}` and `DELETE /conversations/{id}`. PASS on 403/404; FAIL HIGH (read) or CRITICAL (delete) on 2xx. 6 probes total.
12	Privilege escalation	COVERED	`auth/test_forged_groups.py` + `auth/test_cross_persona.py`. Found CRITICAL today (CISO claim forgery is accepted).
13	Forced browsing	COVERED	Block C: `auth/test_forced_browsing.py` — 20-path dictionary of common admin / debug / framework leak paths probed unauthenticated. PASS on 401/403/404; FAIL HIGH on non-trivial 200; `.well-known/security.txt` special-cased as PASS on 200.
14	Missing authorization on APIs	COVERED	Auth layer's negative tests confirmed today that 9 CISO-only routes accept SOC/GRC/Employee tokens (HIGH finding).
15	Server-Side Request Forgery (SSRF)	COVERED	Block A: `fuzz/corpus/ssrf.json` (7 entries: IMDS, loopback, IPv6 loopback, RFC1918, gopher://, file://). Probes every URL-shaped field across all routes.
16	Path traversal	COVERED	`fuzz/corpus/path_traversal.json` (6 entries with `..`, `%2e%2e`, etc.) on every route including those with path params.
17	Credential stuffing / brute force	COVERED	Block C: `auth/test_brute_force.py` — 10 rapid InitiateAuth calls against a UUID-suffixed (non-existent) username. PASS if Cognito returns a throttle code (`LimitExceededException` / `TooManyRequestsException` / `ThrottlingException`) within the window; FAIL HIGH on no throttle. Two rows: `throttle-kicks-in` + `lockout-after-K`.
18	Weak password storage	PARTIAL	Block C: `auth/test_pool_config.py` — 8 black-box config assertions via `cognito-idp:describe_user_pool` (MinimumLength ≥ 12, RequireUppercase/Lowercase/Numbers/Symbols, TemporaryPasswordValidityDays ≤ 7, AccountRecoverySetting present, AllowAdminCreateUserOnly). Storage itself (SRP, hashing) stays out of scope — Cognito's server-side responsibility.

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
19	Session fixation / hijacking	COVERED	Block C: `auth/test_session_swap.py` — (a) CISO creates session X; SOC POSTs to `/chat` with the same `session_id`. FAIL HIGH if SOC's message lands in CISO's history. (b) Stale-token probe — reuse pre-logout IdToken after simulated logout. 2xx recorded as `documented_unsafe` per AC11.
20	JWT vulnerabilities	COVERED	`auth/test_chat_no_signature.py` (alg-none-equivalent), `auth/test_forged_groups.py` (signature manipulation), `auth/test_expired_token.py` (exp validation), `auth/test_token_replay.py` (access-token-instead-of-id). Found CRITICAL + HIGH today.
21	Insecure password reset	COVERED	Block C: `auth/test_password_reset.py` — (a) enumeration: ForgotPassword for known + unknown username; FAIL MEDIUM if outcomes differ. (b) rate-limit: 5 rapid calls against a synthetic username; FAIL MEDIUM if no throttle.
22	Missing or bypassable MFA	OUT OF SCOPE	Per CLAUDE.md: "MFA off" in this demo. Documented intentional gap. If MFA is enabled later, add a probe that attempts to skip the SMS_MFA challenge.

3. Crypto & Data Handling 3 COVERED · 0 PARTIAL · 2 MISSING · 1 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
23	Plaintext transmission of sensitive data	COVERED	Block B: `headers/test_https_only.py` — for every public host (CloudFront SPA, API, Lambda Function URL), GET on `http://` must redirect to `https://` (FAIL HIGH on a 200 or http-to-http redirect) and GET on `https://` must set HSTS with `max-age >= 1 year` (FAIL MEDIUM if missing).
24	Weak / outdated cryptographic algorithms	COVERED	Block B: `headers/test_tls_ciphers.py` — stdlib `ssl` + `socket` only (no `nmap` / `testssl.sh` subprocess). Forces TLS 1.0 / 1.1 handshakes against CloudFront (FAIL HIGH if either succeeds), then negotiates TLS 1.2 and verifies the cipher name is not in the weak-token list (RC4 / 3DES / NULL / EXPORT / anonymous DH / MD5-based).
25	Hardcoded or poorly managed keys	COVERED	Block D: `e2e/tests/bundle-secrets.spec.js::e2e.bundle.hardcoded-keys` — fetches the SPA root HTML, extracts every ` <script src="">` URL, fetches each JS bundle, and regex-scans for AWS access key IDs (`AKIA[0-9A-Z]{16}`), AWS secret-assignment shapes, JWT-shape tokens (medium / flag for review), Slack tokens (`xox[bpa]-...`), GitHub PATs (`ghp_{36}`), and OpenAI/Anthropic key shapes (`sk-...{40}`). PASS on zero hits; FAIL HIGH on AWS/Slack/GitHub/Anthropic-shape; FAIL MEDIUM on JWT-shape. Classifier pinned both in JS (`e2e/lib/bundle-scanner.js`) and in Python parity tests (`tests/test_block_d_bundle_scanner.py`).</td></tr> </table> </div> <div data-bbox="289 952 706 967" data-label="Page-Footer">ARBITER Security Compliance Coverage Report · page 6 of 18</div></script>

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
26	Weak randomness	MISSING	No probe that watches for predictable session_ids or short-cycle UUIDs. ARBITER uses <code>crypto.randomUUID()</code> which is fine, but a regression to <code>Math.random()</code> wouldn't be caught.
27	Improper certificate validation	OUT OF SCOPE	This applies to client-side code. The deployed SPA uses the browser's cert validation; ARBITER doesn't ship a desktop/mobile client.
28	Unnecessary retention of sensitive data	MISSING	No probe of DDB TTL / retention policies. Spec'd in <code>`04-storage.yaml`</code> but not enforced by harness.

4. Configuration & Infrastructure 5 COVERED · 1 PARTIAL · 1 MISSING · 3 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
29	Default credentials	COVERED	Block A: <code>`auth/test_default_creds.py`</code> (5 pairs: admin/admin, admin/password, test/test, arbiter/arbiter, demo/demo123). PASS = Cognito returns <code>NotAuthorizedException</code> ; FAIL severity HIGH if any unexpectedly authenticates.
30	Verbose errors / stack trace leakage	COVERED	Fuzz layer's stack-trace marker check catches <code>`Traceback`</code> , <code>`boto3`</code> , <code>`ClientError`</code> , <code>`aws_request_id`</code> in any response body. Note: earlier review flagged that <code>`"boto3"`</code> is too broad — fixed already.
31	Missing security headers	COVERED	Block B: <code>`headers/test_security_headers.py`</code> (parameterized across SPA root + 3 API routes). Per-header rows: <code>`headers.csp.*`</code> (FAIL MEDIUM if missing, LOW if <code>`unsafe-eval`</code> / <code>`unsafe-inline`</code> / <code>`data:`</code> in script-src), <code>`headers.xfo.*`</code> (clickjacking — FAIL MEDIUM if neither XFO nor CSP <code>`frame-ancestors`</code> restricts framing), <code>`headers.xcto.*`</code> (FAIL LOW if missing / wrong value), <code>`headers.referrer.*`</code> (FAIL LOW if missing or <code>`unsafe-url`</code>). HSTS covered separately by <code>`test_https_only.py`</code> .
32	Open cloud storage / over-permissive IAM	PARTIAL	Manifest-driven static check covers IAM in templates indirectly. No probe attempts unauthenticated GET on S3 bucket URLs or non-CISO access to KMS keys.
33	Exposed admin consoles or sensitive directories	COVERED	Block C: <code>`auth/test_forced_browsing.py`</code> — the 20-path dictionary includes <code>`/admin`</code> , <code>`/console`</code> , <code>`/.git/config`</code> , <code>`/.env`</code> , <code>`/.env.local`</code> , <code>`/swagger`</code> , <code>`/swagger.json`</code> , <code>`/api-docs`</code> , <code>`/openapi.json`</code> , <code>`/wp-admin/`</code> , <code>`/phpmyadmin/`</code> , <code>`/actuator`</code> , <code>`/debug/pprof`</code> , <code>`/metrics`</code> . Same classifier as #13.
34	Directory listing / debug mode in production	MISSING	No probe of CloudFront for directory indices, or of API for debug endpoints (<code>`/debug`</code> , <code>`?debug=1`</code>).

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
35	Misconfigured CORS	COVERED	Block B: `headers/test_cors.py` (route × attacker_origin parametrize). Sends OPTIONS preflights with `Origin: https://evil.com`, `Origin: null`, and `Origin: file://` against 5 representative API routes (15 cells). FAIL HIGH if `Access-Control-Allow-Origin: *` + `Allow-Credentials: true`, or if ACAO echoes the attacker origin back. PASS on no ACAO / fixed-allowed-origin / wildcard-without-credentials.
36	Vulnerable / outdated dependencies	OUT OF SCOPE	GitHub Dependabot is on (we saw the warning on push: 1 critical, 5 moderate). Pin SCA there. The harness doesn't replicate this.
37	Dependency confusion	OUT OF SCOPE	Best handled in build pipeline + private registry config, not runtime probe.
38	Typosquatting	OUT OF SCOPE	Static analysis of `package.json` / `requirements.txt`. Tools like Socket.dev.

5. Build / Supply Chain 2 COVERED · 0 PARTIAL · 0 MISSING · 2 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
39	Compromised build pipelines	OUT OF SCOPE	Requires CI hardening (signed commits, OIDC role assumption, branch protection). Not testable from a black-box runtime harness.
40	Malicious maintainer / package takeover	OUT OF SCOPE	SCA + signed releases.
41	Compromised third-party JS / Magecart skimming	COVERED	Block D: `e2e/tests/bundle-secrets.spec.js::e2e.bundle.sri-on-third-party` — parses every ` <script src="">` and `<link rel="stylesheet" href>` in the SPA root HTML, classifies same-origin vs cross-origin via URL hostname, and asserts every cross-origin tag carries `integrity="..."`. PASS if all third-party tags have SRI; FAIL MEDIUM per missing tag; SKIP with reason "no third-party assets found" if everything is first-party.</td></tr> <tr> <td>42</td><td>Sensitive data in client storage / source maps / comments</td><td>COVERED</td><td>Block D: two probes in `e2e/tests/bundle-secrets.spec.js`. (a) `e2e.bundle.source-maps-in-prod` — HEADs every `<script>.map` URL; FAIL MEDIUM on any 200. (b) `e2e.bundle.sensitive-comments` — regex sweeps the HTML + every JS bundle for `TODO.*(password\`</td></tr> </table> </div> <div data-bbox="289 951 706 967" data-label="Page-Footer">ARBITER Security Compliance Coverage Report · page 8 of 18</div></script>

6. Errors & State 5 COVERED · 0 PARTIAL · 0 MISSING · 0 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
43	Fail-open logic	COVERED	Block H: `fault/test_fail_closed.py` — 5 scenarios (`corrupted-jwt-middle-byte`, `invalid-json-payload`, `no-authorization-header`, `empty-authorization-value`, `non-bearer-scheme`) against a CISO-only route (`GET /token-usage`). PASS on 401/403/400 (fail-closed); FAIL HIGH on 2xx (fail-open: corrupt request accepted); FAIL MEDIUM on 5xx (auth path crashed on bad input).
44	Unhandled exceptions	COVERED	Block H: `fault/test_error_propagation.py` extends the fuzz layer's "no 500s" assertion by triggering known-bad conditions (missing DDB record, cross-pool JWT, oversized prompt) and verifying both a structured response AND a CloudWatch ERROR log line. Plus the fuzz layer's existing 5xx detection still runs.
45	Swallowed errors hiding attacks	COVERED	Block H: `fault/test_error_propagation.py` — 3 scenarios paired with a `.cloudwatch-logged` sub-check per scenario. For each, capture epoch, trigger known-bad condition, scan api_handler CloudWatch log group within a 60 s window for an ERROR-level line containing a scenario-specific needle. PASS if API returned a structured error AND CloudWatch logged it; FAIL LOW if client got an error but ops can't see it (silent error hiding).
46	Race conditions / TOCTOU	COVERED	Block F: `logic/test_race_conditions.py` — (a) 5 concurrent `POST /actions/{id}/approve` against the same action; PASS on exactly 1x2xx + 4x4xx, FAIL HIGH on multiple winners (race window), FAIL MEDIUM on all-5xx (crash). (b) 3 concurrent `DELETE /conversations/{id}` against a freshly-minted CISO conversation; same single-winner contract.
47	Inconsistent state after partial failure	COVERED	Block H: `fault/test_partial_failure_consistency.py` — 3 scenarios. (a) `approve-abort-client`: POST /approve with a 100 ms read timeout, then re-read the action; PASS on consistent state (only-approved or untouched), FAIL HIGH on mixed terminal state. (b) `approve-vs-reject-race`: fire approve + reject in parallel with a thread barrier, then re-read; same single-winner contract. (c) `concurrent-upload-then-scan`: POST /uploads/presign then immediately POST /scan referencing the upload before it could be finalized; PASS on clean refusal or success, FAIL MEDIUM on hang or malformed body.

7. API Top-10 Specific 6 COVERED · 1 PARTIAL · 0 MISSING · 1 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
48	Broken Object-Level Authorization (BOLA)	COVERED	Block C: same surface as IDOR (#11). `auth/test_idor.py` covers per-object cross-persona reads and deletes on `/conversations/{session_id}`.

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
49	Mass assignment	PARTIAL	Block A: <code>`fuzz/corpus/mass_assignment.json`</code> (6 entries: <code>is_admin</code> , <code>persona</code> , <code>cognito:groups</code> , <code>role</code> , <code>user_id</code> , <code>approved</code>) injected as the primary writable field's value on every route. Multi-key body extension (extra top-level keys in the same JSON body) is documented future enhancement — see <code>notes/mass-assignment-extra-keys.md</code> .
50	Excessive data exposure	COVERED	Block F: <code>`logic/test_field_exposure.py`</code> — for each persona × each manifest GET route (skipping <code>`/health`</code> and <code>`{path-param}`</code> routes), fetch the response and walk JSON to depth 6 for sensitive-field patterns. FAIL HIGH on <code>`password`</code> / <code>`password_hash`</code> / <code>`secret`</code> / <code>`api_key`</code> / <code>`private_key`</code> / <code>`aws_access_key`</code> shapes; FAIL MEDIUM on cross-persona <code>`cognito:groups`</code> or cross-user <code>`email`</code> ; FAIL LOW on <code>`_internal`</code> / <code>`internal_id`</code> / Mongo <code>`_id`</code> / <code>`__v`</code> . 38 (persona × route) cases parameterised.
51	Lack of rate limiting / resource consumption	COVERED	Block E: <code>`dos/test_rate_limit.py`</code> — for 5 representative routes (<code>`get-findings`</code> , <code>`get-conversations`</code> , <code>`get-dashboard`</code> , <code>`get-agent-status`</code> , <code>`post-chat`</code>), sends a sustained burst at <code>`--dos-rps`</code> (default 20, hard ceiling 100) for <code>`--dos-duration-seconds`</code> (default 5, hard ceiling 30). PASS on ≥1 429; FAIL MEDIUM if no 429 and latency flat (rate limiting absent); FAIL HIGH if any 500 / transport drop / monotonic latency growth (API buckling).
52	Broken function-level authorization (BFLA)	COVERED	Same as #14 — auth layer's per-route persona tests cover BFLA. Confirmed broken today.
53	Unsafe consumption of third-party APIs	COVERED	Block H: <code>`fault/test_unsafe_third_party.py`</code> — 2 scenarios on the specialist surface. (a) <code>`jira-error-leak`</code> : POST <code>/chat</code> with a crafted prompt aimed at the jira specialist tool ("ticket ID -1 -- DROP TABLE issues"); verify the master orchestrator doesn't leak a raw stack trace from the specialist back through the chat surface. (b) <code>`specialist-latency`</code> : POST <code>/chat</code> with a multi-tool prompt and a 30 s hang threshold; PASS on clean completion within the cap, FAIL MEDIUM on hang (transport drop) or 5xx (orchestrator crash).
54	GraphQL-specific abuse	OUT OF SCOPE	ARBITER doesn't expose GraphQL.
55	Cross-Site Request Forgery (CSRF)	COVERED	Block B: <code>`headers/test_csrf.py`</code> . For every POST / PUT / PATCH / DELETE route in the manifest, fire the request with NO <code>`Authorization:`</code> header but WITH a fake <code>`Cookie: arbiter.tokens=...`</code> . Expected: 401 / 403. FAIL HIGH if the API returns 2xx — that would mean a cookie-based auth fallback exists, exposing the surface to classic CSRF.

8. Client-Side 5 COVERED · 0 PARTIAL · 0 MISSING · 0 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
56	Clickjacking	COVERED	Block B (two-pronged): (a) <code>`headers/test_clickjacking.py::test_iframe_embed_blocked_via_headers`</code> — header-side check that XFO DENY/SAMEORIGIN or CSP <code>`frame-ancestors`</code> restricts framing (FAIL MEDIUM if neither); (b) <code>`e2e/tests/clickjacking.spec.js`</code> — real-browser Playwright spec that wraps the SPA URL in an <code><iframe></code> and asserts the dashboard does not render inside it.
57	Open redirects	COVERED	Block A: <code>`fuzz/corpus/open_redirects.json`</code> (7 entries: bare host, protocol-relative, full https, UNC, mixed slash, subdomain confusion, <code>`@`</code> userinfo bypass) + dedicated probe <code>`fuzz/test_open_redirects.py`</code> against the Cognito Hosted UI <code>`/login?redirect_uri=`</code> with Location-header inspection. ARBITER's own API has no redirect routes — the Hosted UI is the exposure surface.
58	Prototype pollution	COVERED	Block A: <code>`fuzz/corpus/prototype_pollution.json`</code> (5 entries: <code>`__proto__.isAdmin`</code> , <code>`constructor.prototype`</code> , <code>`__proto__.toString`</code> , nested, array-typed). Lambda runtime is Python — these are regression detectors for a future Node Lambda.
59	Tabnabbing	COVERED	Block D: <code>`e2e/tests/bundle-tabnabbing.spec.js`</code> — CISO-authenticated sweep over Dashboard / Settings / Integrations. For every <code></code> whose href hostname differs from the SPA host, asserts the <code>`rel`</code> attribute contains BOTH <code>`noopener`</code> AND <code>`norereferrer`</code> . PASS on protected; FAIL MEDIUM per link missing either token; SKIP on pages with no external <code>`target="_blank"`</code> links.
60	Sensitive data in client storage / source maps / comments	COVERED	Same as #42 (Block D — source-maps + sensitive-comments probes).

9. Logic & Workflow 1 COVERED · 0 PARTIAL · 0 MISSING · 2 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
61	Workflow bypass	COVERED	Block F: `logic/test_action_state_machine.py` — 4 probes against the action lifecycle: (a) `skip-approve` — `POST /actions/{id}/execute` without prior approval; (b) `double-approve` — same approver twice; (c) `reject-after-execute` — reject a COMPLETED / terminal action; (d) `escalate-from-terminal`. PASS on 400/403/404/409/422 refusal; FAIL HIGH on 2xx for skip-approve (direct workflow bypass) or any 5xx; FAIL MEDIUM on 2xx for the others (idempotency / lifecycle break).
62	Price / quantity manipulation	OUT OF SCOPE	Not applicable to a compliance dashboard.
63	Abuse of trial, referral, or exempt flows	OUT OF SCOPE	Not applicable.

10. Denial of Service 2 COVERED · 0 PARTIAL · 0 MISSING · 1 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
64	Application-layer DoS	COVERED	Block E: `dos/test_payload_oversize.py` — for 3 JSON-accepting routes (`post-chat`, `post-jira-tickets`, `post-actions`), sends a 32-byte baseline, then 1 MB, then 10 MB JSON bodies. PASS if the API rejects with 4xx OR if 10 MB latency stays below 3× baseline. FAIL MEDIUM if 10 MB latency is 3–5× baseline (degraded but absorbed). FAIL HIGH if 10 MB latency > 5× baseline, the call timed out (> 30 s), or the server returned 5xx. Augments the LLM cost-DoS probe at the API layer.
65	Resource exhaustion	COVERED	Block E: `dos/test_resource_exhaustion.py` — fans out 10 concurrent requests against `POST /scan` (or `POST /chat` under `--include-bedrock-dos`) using `concurrent.futures.ThreadPoolExecutor(max_workers=10)`. PASS on a clean mix of 200/429/503; FAIL HIGH on any 500 (server crash) or zero-status (transport drop); FAIL MEDIUM on any malformed response (truncated JSON / wrong content-type) or unexpected 4xx. Plus the cost-DoS probe still covers the LLM-spend exhaustion vector.
66	Network-layer DDoS	OUT OF SCOPE	Would actually attack the dev env. CloudFront has AWS Shield Standard by default; verification is via AWS-side metrics.

11. Logging & Monitoring 3 COVERED · 0 PARTIAL · 0 MISSING · 3 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
67	Insufficient logging of security events	COVERED	Block G: `logging_audit/test_security_events_logged.py` — 4 scenarios (`forged-token`, `cross-persona`, `legitimate-approve`, `brute-force`). For each, capture epoch, trigger event, sleep 3 s, scan `dev-st21arbiter-poc-audit-log` DDB table with a FilterExpression on `timestamp >= start_iso` plus client-side needle match. PASS on ≥ 1 matching row; FAIL HIGH on zero matches.
68	Logging sensitive data	COVERED	Block G: `logging_audit/test_log_redaction.py` — 3 scenarios. (a) JWT not logged: send IdToken in Authorization, search CloudWatch via `logs:FilterLogEvents` for first + last 40 char fragments. (b) Body field not logged: POST `/chat` with `harness-canary-<uuid>-secret`, search CloudWatch. (c) Email not logged in errors: trigger Cognito InitiateAuth failure with synthetic email, search CloudWatch. PASS on zero canary hits; FAIL HIGH per hit.
69	No alerting / monitoring	OUT OF SCOPE	CloudWatch Alarms / EventBridge rule audit. Separate from runtime probe.
70	Mutable or attacker-deletable logs	OUT OF SCOPE	CloudWatch immutability / S3 Object Lock configuration audit.
71	Log injection	COVERED	API-layer reflection check covered by Block A (see #10). Block G: `logging_audit/test_log_injection_downstream.py` — for each payload in `fuzz/corpus/log_injection.json`, wraps the payload with a UUID marker, POSTs as the `/chat` prompt, then reads CloudWatch and counts matching events. PASS on 1 event (no split) with no ANSI residue; FAIL MEDIUM if CR/LF split the line into >1 events (log forgery); FAIL LOW on single-line but ANSI escape present; PASS on 0 events (API didn't log; not the classifier's finding).
72	No incident response plan	OUT OF SCOPE	Process / runbook, not a runtime test.

12. LLM Top-10 4 COVERED · 0 PARTIAL · 0 MISSING · 2 OUT OF SCOPE

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
73	Prompt injection	COVERED	LLM red-team: 20 curated jailbreaks + 10 generative transformations. Found today that 26 of 30 were correctly refused (Bedrock Guardrails) but classifier missed the phrasing (harness fix planned).

#	ITEM	STATUS	WHERE IT IS COVERED / GAP
74	Insecure output handling	COVERED	Block H: <code>`fault/test_unsafe_third_party.py`</code> — 2 scenarios on the API boundary. (a) <code>`xss-in-json`</code> : POST /chat asking the model to emit a verbatim <code>`<script>`</code> payload; PASS if the response is valid <code>`application/json`</code> (so the SPA can <code>`textContent`</code> it safely), FAIL MEDIUM if the response has a non-JSON content-type and the raw payload made it through. (b) <code>`unverified-link-suggestion`</code> : POST /chat asking the model to suggest a URL; PASS if the URL is absent or appears as plain text, FAIL LOW if the response wraps the URL in an <code>`<a href>`</code> tag (phishing affordance). The browser-side DOM-rendering check is documented in <code>`e2e/tests/llm-output-xss.spec.js`</code> as the SPA-side counterpart (planned).
75	Training-data poisoning	OUT OF SCOPE	Requires access to Bedrock model training pipeline. Not testable for a hosted model.
76	Sensitive information disclosure	COVERED	<code>`exfil.list-other-ciso-conversations`</code> , <code>`exfil.cross-tenant-sharepoint`</code> , <code>`exfil.knowledge-base-dump`</code> probes. Reinforced by jira black-box check (AC19) for 12-digit AWS account IDs in responses.
77	Excessive agency	COVERED	<code>`tool-abuse.sharepoint-path-traversal`</code> , <code>`tool-abuse.awsconfig-cross-account`</code> , <code>`tool-abuse.zscaler-crafted-url`</code> probes.
78	Model / plugin supply-chain risks	OUT OF SCOPE	Bedrock model provenance + AgentCore container audit. Separate AWS-side audit.

Out-of-scope appendix

The 16 items below are real items on the compliance checklist. They are not testable by an automated runtime harness against a deployed app, so they are routed to their proper control: a different tool (SCA, AWS config audit), a different team (incident response runbook review), or a different process control (CI hardening, signed releases). They are listed here deliberately so an auditor can verify nothing was dropped.

#	ITEM	CATEGORY	WHERE IT IS ACTUALLY COVERED
22	Missing or bypassable MFA	<i>2. Authentication & Access Control</i>	Per CLAUDE.md: "MFA off" in this demo. Documented intentional gap. If MFA is enabled later, add a probe that attempts to skip the SMS_MFA challenge.
27	Improper certificate validation	<i>3. Crypto & Data Handling</i>	This applies to client-side code. The deployed SPA uses the browser's cert validation; ARBITER doesn't ship a desktop/mobile client.
36	Vulnerable / outdated dependencies	<i>4. Configuration & Infrastructure</i>	GitHub Dependabot is on (we saw the warning on push: 1 critical, 5 moderate). Pin SCA there. The harness doesn't replicate this.
37	Dependency confusion	<i>4. Configuration & Infrastructure</i>	Best handled in build pipeline + private registry config, not runtime probe.
38	Typosquatting	<i>4. Configuration & Infrastructure</i>	Static analysis of `package.json` / `requirements.txt`. Tools like Socket.dev.
39	Compromised build pipelines	<i>5. Build / Supply Chain</i>	Requires CI hardening (signed commits, OIDC role assumption, branch protection). Not testable from a black-box runtime harness.
40	Malicious maintainer / package takeover	<i>5. Build / Supply Chain</i>	SCA + signed releases.
54	GraphQL-specific abuse	<i>7. API Top-10 Specific</i>	ARBITER doesn't expose GraphQL.
62	Price / quantity manipulation	<i>9. Logic & Workflow</i>	Not applicable to a compliance dashboard.
63	Abuse of trial, referral, or exempt flows	<i>9. Logic & Workflow</i>	Not applicable.

#	ITEM	CATEGORY	WHERE IT IS ACTUALLY COVERED
66	Network-layer DDoS	10. Denial of Service	Would actually attack the dev env. CloudFront has AWS Shield Standard by default; verification is via AWS-side metrics.
69	No alerting / monitoring	11. Logging & Monitoring	CloudWatch Alarms / EventBridge rule audit. Separate from runtime probe.
70	Mutable or attacker-deletable logs	11. Logging & Monitoring	CloudWatch immutability / S3 Object Lock configuration audit.
72	No incident response plan	11. Logging & Monitoring	Process / runbook, not a runtime test.
75	Training-data poisoning	12. LLM Top-10	Requires access to Bedrock model training pipeline. Not testable for a hosted model.
78	Model / plugin supply-chain risks	12. LLM Top-10	Bedrock model provenance + AgentCore container audit. Separate AWS-side audit.

Methodology & reproduction

Harness shape

The harness is a black-box test suite that runs against the deployed dev environment. It does not have access to backend source at runtime — every assertion is made on what the deployed API, the deployed SPA, or the deployed AgentCore runtimes actually return. A nightly cost cap (default \$1.00) gates LLM-spending probes so a runaway fan-out cannot exhaust the demo budget. Layers that touch the LLM surface (`llm`, `fault`) account against that cap; the other seven layers are zero-cost.

The nine layers

LAYER	WHAT IT TESTS
<code>e2e</code>	Playwright sweeps the SPA per persona and asserts page-level access.
<code>fuzz</code>	10 curated payload corpora cross-product over every API route.
<code>auth</code>	JWT forgery, IDOR, brute-force, session fixation, password reset.
<code>headers</code>	HTTPS / HSTS / TLS / CSP / CORS / CSRF / clickjacking.
<code>llm</code>	20 jailbreaks + 10 generative red-team prompts against the master.
<code>dos</code>	Rate limiting, oversized payloads, concurrent fan-out.
<code>logic</code>	Workflow bypass, race conditions, sensitive-field exposure.
<code>logging_audit</code>	Audit-log writes + CloudWatch redaction + log injection.
<code>fault</code>	Fail-closed, error propagation, partial failure, LLM output safety.

Run cadence

The full harness runs on a daily cadence against the dev environment. Each run produces a `report.html` + `summary.md` + `report.json` in `tests-adversarial/test-reports/<run-id>/`, and the dated PDF in `docs/adversarial_run_findings_<date>.pdf`. After every "build block" lands (one new test layer or a refresh of an existing one), this compliance matrix is regenerated so the coverage percentage moves with reality.

Reproducing this report

From `tests-adversarial/`:

```
npm run test:all
```

 — run the nine layers and aggregate.

```
python3.13 -m scripts.compliance_report
```

 — regenerate this PDF from

```
docs/security_compliance_coverage.md
```

.

This document is generated from `docs/security_compliance_coverage.md`. The markdown is the single source of truth; regenerate this PDF after every build block lands so the coverage percentage moves with the harness. Last generated 2026-06-10.

Generated 2026-06-10 · harness rev `c9498c49` · source `docs/security_compliance_coverage.md` · re-run with
`python3.13 -m scripts.compliance_report`.